



Progress Report

By Wayne, Young Joo, and Kelson

Outline

- Introduction to Deep Fashion
 - Introduction
 - Goal
- Object Detection API
 - Introduction
 - Young Joo
 - Kelson
 - Wayne
 - Problems & Conclusion
- Bad Data
 - Data cleaning
- Future Plan

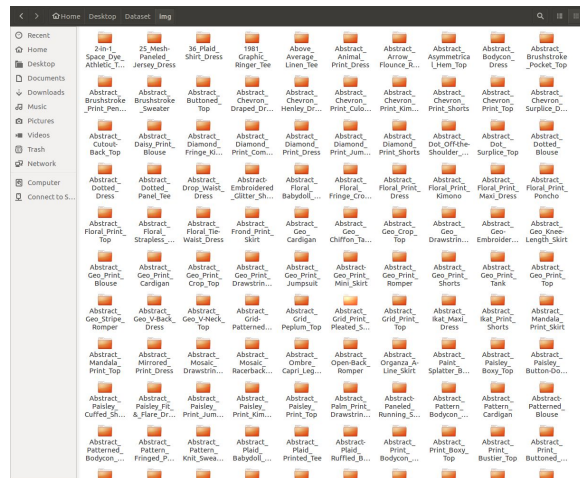


Deep Fashion - Introduction

- A large scale dataset
 - Chinese University of Hong Kong(CHUK)
 - 289,222 clothes
 - 50 broad & 5612 specific categories
 - Bounding box & clothing type/label annotations

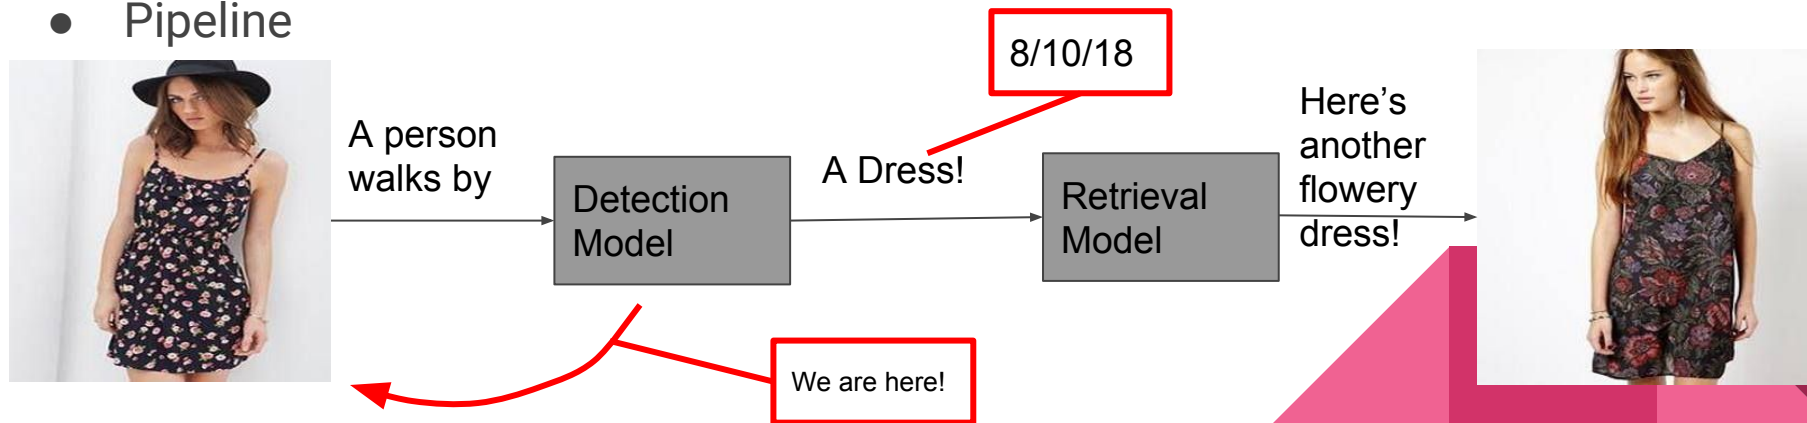


Flower & Dot Cami Dress



Deep Fashion - Goal

- In the far future...
 - Recommend clothing to customers that walk by
 - Based on the clothes in the data, recommend those of the similar styles
- Before August 10th
 - Detect the clothes the customer is wearing
- Pipeline

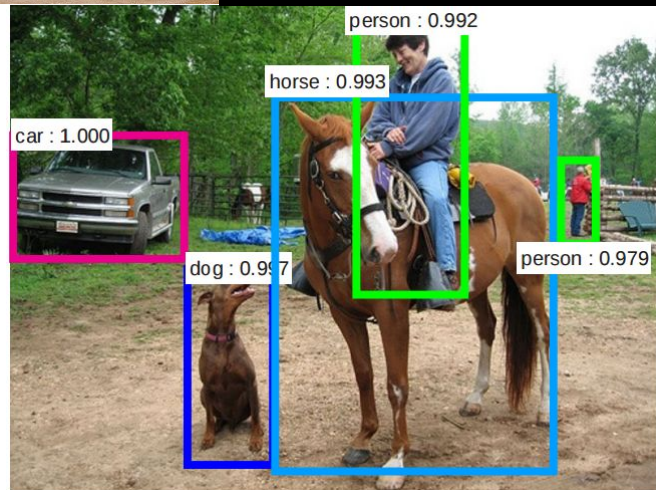




Object Detection API

Our Models

Object Detection


- An API provided by GOOGLE
 - Provides model zoo with pretrained ckpt
 - Pretrained on datasets including
 - COCO (objects)
 - Kitti (street view)
 - OpenImages (objects)
 - AVA (human actions e.g. listening, holding)
- How to use
 - Prepare data
 - Select proper model
 - Finetune



Young-Joo

- Ssd_mobilenet_v2 (pre-trained on COCO)
- 50 categories
- Trained for 4 days
 - Instead of converging, loss fluctuated
 - Learning rate also fluctuated
- Performed ok, but far from well

Losses/TotalLoss



LearningRate/LearningRate/learning_rate




image-0
step 169,664
Fri Jul 13 2018 17:41:14 GMT+0800 (CST)



image-1
step 169,664
Fri Jul 13 2018 17:41:14 GMT+0800 (CST)




image-6
step 110,994
Fri Jul 13 2018 10:00:22 GMT+0800 (CST)

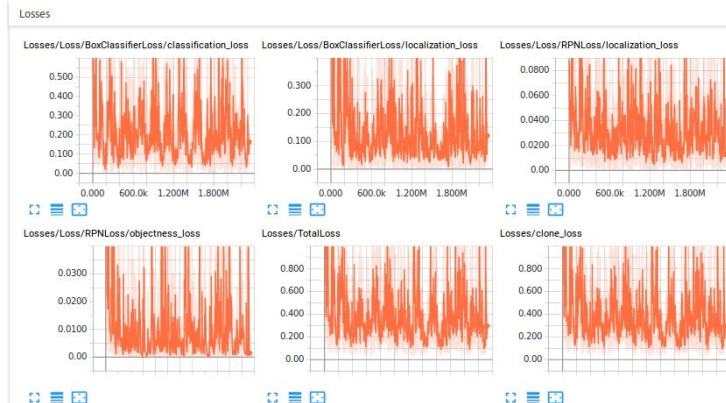
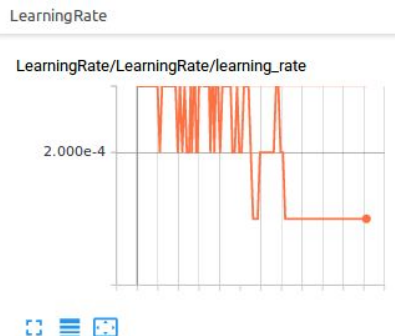


image-1
step 152,598
Fri Jul 13 2018 15:44:35 GMT+0800 (CST)




Kelson

- Model: faster_rcnn_inception_v2
 - Pretrained on COCO
- Number of Categories: 50
- Training time:
 - 2d 23h 8m 0s
 - 2,318,309 steps
- Fluctuations
 - Learning Rate
 - Losses
- Poor performance
 - Everything is jumpsuit?



Wayne

- Model
 - Faster RCNN Inception Resnet v2 atrous
 - Pretrained on COCO
- Reasoning
 - This model has ckpts pretrained on both COCO and OpenImage
 - Two datasets are both more relevant to our dataset
 - Can see if it is the model structure or the pretraining dataset that matters
- 3 categories instead of 50
 - Trained for (196.5k steps) 2d15h39m11s
 - Model had great test mAP(0.9962) when at 99.72k steps
 - Works well with 3 categories if only there were no Overfitting
 - Too few categories compared with the complexity/capacity of the model
 - No Dropout layer
 - No Early stopping

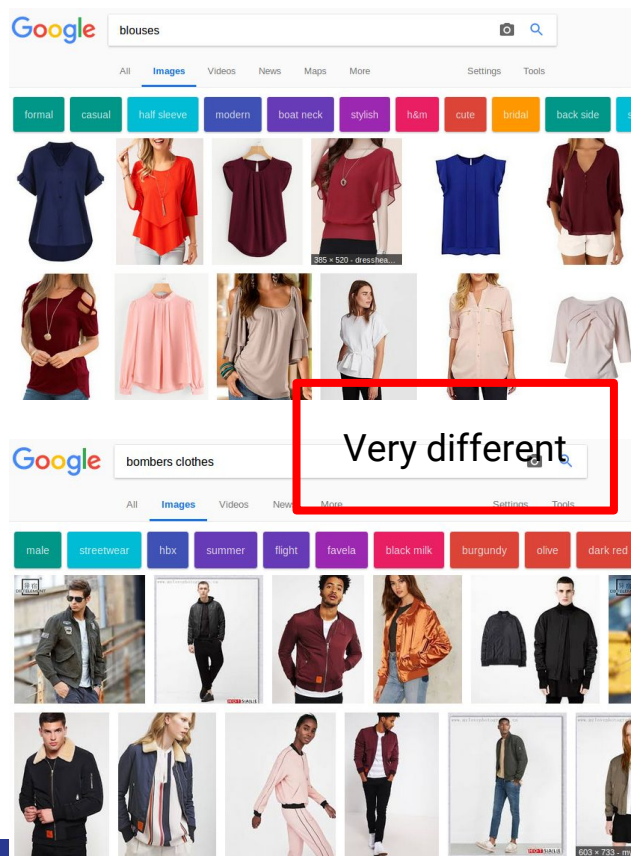




Problems We Faced

Problems, Causes, and Solutions

Problem 1: All blouses are bombers...?



Right next to
each other...?

| category_name | category_type |
|---------------|---------------|
| Anorak | 1 |
| Blazer | 1 |
| Blouse | 1 |
| Bomber | 1 |
| Button-Down | 1 |
| Cardigan | 1 |
| Flannel | 1 |
| Halter | 1 |
| Henley | 1 |
| Hoodie | 1 |
| Jacket | 1 |
| Jersey | 1 |

Cause: Mismatch between labels and indices

```
cats = ['Anorak', 'Blazer', 'Blouse', 'Bomber', 'Button-Down', 'Cardigan', ...]  
label_num = 1 to 50  
cats[50] = out_of_bounds_ERROR!
```

```
def save_to_record(images_addr, labels, files_name, bboxes, cats, option = None):  
    if option == None:  
        filename = TFRERCORD_NAME  
    else:  
        filename = option  
    writer = tf.python_io.TFRecordWriter(filename)  
    length = len(images_addr)  
    count = 0  
    for i in range(length):  
        if count % 1000 == 0:  
            print("%d out of %d saved" % (count, length))  
        image = cv.imread(PATH_TO_FILES + images_addr[i])  
        image = cv.cvtColor(image, cv.COLOR_BGR2RGB)  
        (h, w) = image.shape[:2]  
        with tf.gfile.GFile(PATH_TO_FILES + images_addr[i] + 'resized', 'rb') as fid:  
            encoded_jpg = fid.read()  
  
        label_num = labels[i]  
        example = tf.train.Example(features=tf.train.Features(feature={  
            'image/height': int64_feature(299),  
            'image/width': int64_feature(299),  
            'image/filename': bytes_feature(files_name[i].encode()),  
            'image/source_id': bytes_feature(files_name[i].encode()),  
            'image/encoded': bytes_feature(encoded_jpg),  
            'image/format': bytes_feature(b'jpg'),  
            'image/object/bbox/xmin': float_list_feature([bboxes[0][i] / w]),  
            'image/object/bbox/xmax': float_list_feature([bboxes[1][i] / w]),  
            'image/object/bbox/ymin': float_list_feature([bboxes[2][i] / h]),  
            'image/object/bbox/ymax': float_list_feature([bboxes[3][i] / h]),  
            'image/object/class/text': bytes_feature(cats[label_num].encode()),  
            'image/object/class/label': int64_feature(labels[i])  
        }))  
        writer.write(example.SerializeToString())  
        count = count + 1  
    writer.close()  
    return example
```

labelmap - Notepad

File Edit Format View Help

```
item {  
  id:1  
  name:'Anorak'  
}  
item {  
  id:2  
  name:'Blazer'  
}  
item {  
  id:3  
  name:'Blouse'  
}  
item {  
  id:4  
  name:'Bomber'  
}  
item {  
  id:5  
  name:'Button-Down'  
}  
item {  
  id:6  
  name:'Cardigan'  
}  
item {  
  id:7  
  name:'Flannel'  
}  
item {  
  id:8  
  name:'Halter'  
}  
item {  
  id:9  
  name:'Henley'  
}  
item {  
  id:10  
  name:'Hoodie'  
}
```

Solution: Shift by 1

```
def save_to_record(images_addr, labels, files_name, bboxes, cats, option = None):
    if option == None:
        filename = TFRECORD_NAME
    else:
        filename = option
    writer = tf.python_io.TFRecordWriter(filename)
    length = len(images_addr)
    count = 0
    for i in range(length):
        if count % 1000 == 0:
            print("%d out of %d saved" % (count, length))
        image = cv.imread(PATH_TO_FILES + images_addr[i])
        image = cv.cvtColor(image, cv.COLOR_BGR2RGB)
        (h, w) = image.shape[:2]
        #####Adding this two lines to try
        split = images_addr[i].split(".")
        with tf.gfile.GFile(PATH_TO_FILES + split[0] + "revised.jpg", 'rb') as fid:
            encoded_jpg = fid.read()
        label_num = labels[i]
        example = tf.train.Example(features=tf.train.Features(feature={
            'image/height': int64_feature(299),
            'image/width': int64_feature(299),
            'image/filename': bytes_feature(files_name[i].encode()),
            'image/source_id': bytes_feature(files_name[i].encode()),
            'image/encoded': bytes_feature(encoded_jpg),
            'image/format': bytes_feature(b'jpg'),
            'image/object/bbox/xmin': float_list_feature([bboxes[0][i] / w]),
            'image/object/bbox/xmax': float_list_feature([bboxes[1][i] / w]),
            'image/object/bbox/ymin': float_list_feature([bboxes[2][i] / h]),
            'image/object/bbox/ymax': float_list_feature([bboxes[3][i] / h]),
            'image/object/class/text': bytes_feature(cats[label_num-1].encode()),
            'image/object/class/label': int64_feature(labels[i])
        }))
        writer.write(example.SerializeToString())
        count = count + 1
    writer.close()
    return example
```

Problem 2: Results were Inaccurate

- Fine-tuning pre-trained models
 - can only change the configuration file
 - # of categories
 - the dimensions that the images will be converted to
 - # of maximum training steps
 - Batch size
 - Path to tfrecord files

```
1 # SSD with Mobilenet v2 configuration for MSCOCO Dataset.
2 # Users should configure the fine_tune_checkpoint field in the train_config as
3 # well as the label_map_path and input_path fields in the train_input_reader and
4 # eval_input_reader. Search for "PATH_TO_BE_CONFIGURED" to find the fields that
5 # should be configured.
6
7 model {
8   ssd {
9     num_classes: 90
10    box_coder {
11      faster_rcnn_box_coder {
12        y_scale: 10.0
13        x_scale: 10.0
14        height_scale: 5.0
15        width_scale: 5.0
16      }
17    }
18    matcher {
19      argmax_matcher {
20        matched_threshold: 0.5
21        unmatched_threshold: 0.5
22        ignore_thresholds: false
23        negatives_lower_than_unmatched: true
24        force_match_for_each_row: true
25      }
26    }
27    similarity_calculator {
28      iou_similarity {
29      }
30    }
31    anchor_generator {
32      ssd_anchor_generator {
33        num_layers: 6
34        min_scale: 0.2
35        max_scale: 0.95
36        aspect_ratios: 1.0
37        aspect_ratios: 2.0
38        aspect_ratios: 0.5
39        aspect_ratios: 3.0
40        aspect_ratios: 0.3333
41      }
42    }
43    image_resizer {
44      fixed_shape_resizer {
45        height: 299
46        width: 299
47      }
48    }
49    box_predictor {
50      convolutional_box_predictor {
51        min_depth: 0
52        max_depth: 0
53        num_layers_before_predictor: 0
54        use_dropout: false
55        dropout_keep_probability: 0.0
56        kernel_size: 1
57        box_code_size: 4
58        apply_sigmoid_to_scores: false
59        conv_hypers {
```


Cause: Images of Different Size

- A bug in the tfrecord image size conversion function
 - Need to resize it ourselves



Oleksandr Savsunenko

Follow

Head of AI Lab at Skylum Software, doer/maker/dreamer, father

Jan 24 · 3 min read

How Tensorflow's `tf.image.resize` stole 60 days of my life

That's a short warning to all Tensorflow users working with visual content. Short notice: don't use any `tf.image.resize` functions!

- Resizing images doesn't affect the quality of images
 - The category of each article of clothing - invariant to the size of the image

Solution: Resize them, but how?

- Inside TFRecord

- Complicated due to bounding boxes
- Too time-consuming and inefficient

If you resize it here, you need to read it again

Has to be the original height and width to normalize bounding boxes

```
def save_to_record(images_addr, labels, files_name, bboxes, cats, option = None):
    if option == None:
        filename = TFRECORD_NAME
    else:
        filename = option
    writer = tf.python_io.TFRecordWriter(filename)
    length = len(images_addr)
    count = 0
    for i in range(length):
        if count % 1000 == 0:
            print("%d out of %d saved" % (count, length))
        image = cv.imread(PATH_TO_FILES + images_addr[i])
        image = cv.cvtColor(image, cv.COLOR_BGR2RGB)
        (h, w) = image.shape[:2]
        with tf.gfile.GFile(PATH_TO_FILES+images_addr[i] + 'resized', 'rb') as fid:
            encoded_jpg = fid.read()

        label_num = labels[i]
        example = tf.train.Example(features=tf.train.Features(feature={
            'image/height': int64_feature(299),
            'image/width': int64_feature(299),
            'image/filename': bytes_feature(files_name[i].encode()),
            'image/source id': bytes_feature(files_name[i].encode()),
            'image/encoded': bytes_feature(encoded_jpg),
            'image/format': bytes_feature(b'jpg'),
            'image/object/bbox/xmin': float_list_feature([bboxes[0][i] / w]),
            'image/object/bbox/xmax': float_list_feature([bboxes[1][i] / w]),
            'image/object/bbox/ymin': float_list_feature([bboxes[2][i] / h]),
            'image/object/bbox/ymax': float_list_feature([bboxes[3][i] / h]),
            'image/object/class/text': bytes_feature(cats[label_num].encode()),
            'image/object/class/label': int64_feature(labels[i])
        }))
        writer.write(example.SerializeToString())
        count = count + 1
    writer.close()
    return example
```

Solution: Resize them, but how?

- Preprocess the Images

- Still Complicated due to bounding boxes
- Much more efficient

Record height and width of the original image



Input

Output

```
image = cv.imread(PATH_TO_FILES + images_addr[i])
image = cv.cvtColor(image, cv.COLOR_BGR2RGB)
(h, w) = image.shape[:2]
with tf.gfile.GFile(PATH_TO_FILES+images_addr[i] + 'resized', 'rb') as fid:
    encoded_jpg = fid.read()

label_num = labels[i]
example = tf.train.Example(features=tf.train.Features(feature={
    'image/height': int64_feature(299),
    'image/width': int64_feature(299),
    'image/filename': bytes_feature(files_name[i].encode()),
    'image/source_id': bytes_feature(files_name[i].encode()),
    'image/encoded': bytes_feature(encoded_jpg),
```

convert_size.py

```
img = Image.open(PATH_TO_FILES + images_addr[i])
img = img.resize((299, 299), PIL.Image.ANTIALIAS)
img.save(PATH_TO_FILES + split[0] + "revised.jpg")
```


But save the resized ones

Problem 3: Poor Performance

- Even after four days of training...

Terrible accuracy for most classes

Losses/Loss/classification_loss



Super high

Losses/Loss/localization_loss



Acceptable

```
INFO:root:Writing metrics to tf summary.
INFO:root:Losses/Loss/classification_loss: 6.998816
INFO:root:Losses/Loss/localization_loss: 0.617418
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Anorak: 0.000257
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Blazer: 0.026980
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Bloxxer: 0.077074
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Bomber: 0.008662
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Button-Down: 0.016414
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Caftan: 0.000000
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Cape: nan
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Capris: 0.235925
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Cardigan: 0.437638
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Chinos: 0.187986
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Coat: 0.004308
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Coverup: 0.000000
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/culottes: 0.127387
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Cutoffs: 0.192884
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Dress: 0.298415
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Flannel: 0.126319
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Gauchos: 0.002021
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Halter: 0.000000
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Henley: 0.088897
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Hoodie: 0.418715
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Jacket: 0.412425
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Jean: 0.547041
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Jeggings: 0.076053
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Jersey: 0.021452
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Jodhpurs: 0.002107
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Joggers: 0.306339
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Jumpsuit: 0.114014
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Kaftan: 0.001600
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Kimono: 0.059806
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Leggings: 0.588008
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Nightdress: nan
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Onesie: 0.000000
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Parka: 0.036631
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Peacoat: 0.002233
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Poncho: 0.137744
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Robe: 0.000024
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Romper: 0.023940
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Sarong: 0.000090
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Shirt: nan
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Shorts: 0.471968
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Skirt: 0.436888
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Sundress: nan
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Sweater: 0.465493
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Sweatpants: 0.332838
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Sweatshorts: 0.103870
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Tank: 0.457678
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Tee: 0.654558
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Top: 0.094506
INFO:root:PascalBoxes_PerformanceByCategory/AP@0.5IOU/Turtleneck: 0.045158
INFO:root:PascalBoxes_Precision/mAP@0.5IOU: 0.169396
INFO:root:Metrics written to tf summary.
```


Cause: Terrible Dataset

- Data is important

- According to

<https://petewarden.com/2018/05/28/why-you-need-to-improve-your-training-data-and-how-to-do-it/>


- model with real users early and often. You'll always be able to swap out an improved model down the road, and maybe see better results, but you have to get the data right first. Deep learning still obeys the fundamental computing law of "garbage in, garbage out", so even the best model will be limited by flaws in your training set. By picking a model and testing it, you'll be able to understand what those flaws are and start improving them.

- Model learns garbage from bad data

Why you need to improve your training data, and how to do it

FOLLOW @PETEWARDEN ON TWITTER

Tweets by @petewarden



Bad Data

What We are Working on Right Now

Bad Data

- Inaccurate Data
 - Sorting/Labeling
 - E.g. Suit inside Tees folder
- Poor Quality
 - Bad resolutions
 - Other items covering the object
- Images with:
 - Multiple items
 - (of same object)
 - No models
 - Images with only clothes



Note: These are actual images from our data set

Solution: Manually Clean the Dataset!

- Keep images only if there is one person
 - Remove those with mannequin
 - Remove those with multiple people
 - Multiple articles of clothing
 - Remove those with no person
- Merge or delete categories
 - Some types of clothing are hybrid
 - If it looks like a dress, then it is a dress, but not a romper
 - Anoraks and Bombers are basically Jackets

Anorak



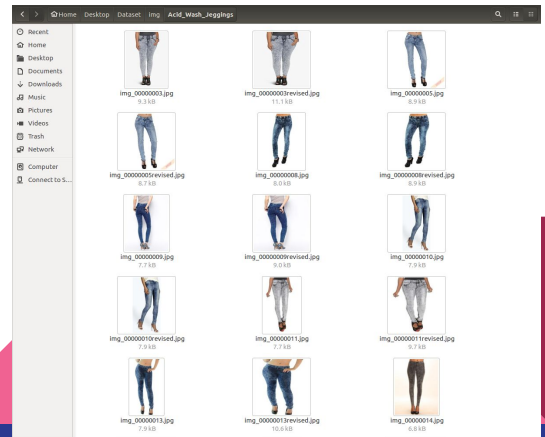
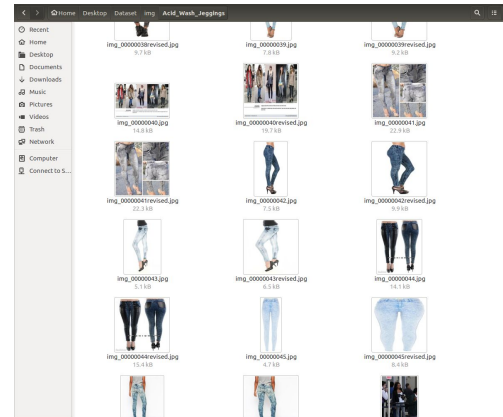
Jacket



Bomber



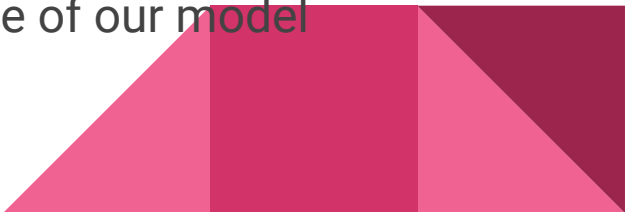
Romper that looks like a dress



Future Plan

What We are Working on Right Now

Future Plan

- Still in progress
 - We each pick 3 categories from either TOP, BOTTOM, FULL-BODY
 - Collect 1k images for each categories
 - If synchronously finished cleaning data:
 - Train with all the cleaned data we have
 - Else:
 - Train our own cleaned data
 - Then, keep on cleaning the data during the training
 - Incrementally clean and add more categories
 - Come up with more ways to improve the performance of our model
- 

Citations

https://cdn-images-1.medium.com/max/655/1*62JsN-0BZ2r11QdhNLSN8A.png

https://cdn-images-1.medium.com/max/603/1*N8H9Z3RDQz2LW_NHh9l3Fw.jpeg

https://i.ytimg.com/vi/_zZe27JYi8Y/hqdefault.jpg

<http://mmlab.ie.cuhk.edu.hk/projects/DeepFashion.html>



Questions? Advice?

你們是老大!!!!!!!!!!!!!!!!!!!!

Fashionable
Machiko Rabbit

